



CHAPTER 6 : Accessing Software



User Environment

HPC systems can support a large and complex range of software

- Operating System (Ubuntu, CentOS)
- Compilers (e.g. Intel, GCC, Cray, PGI)
- Debuggers and Profilers (e.g. MAP , DDT)
- Performance mathematical libraries (e.g. Intel's MKL, Lapack, Petsc)
- Parallel programming libraries (e.g. MPI)
- File format libraries for parallel IO (e.g. HDF5)

Project groups may be expected to manage the installation of their own software stack.

All of the above software is not immediately available as soon as you log in.



Environment Modules

To prevent conflicts between software names and versions, applications and libraries are not installed in the standard directory locations.

Modules modify the environment to easily locate software, libraries, documentation, or particular versions of the software

```
module load openmpi
```



Module Commands

Command	Description
<code>module avail</code>	Show available modules
<code>module list</code>	List loaded modules
<code>module load <i>modulename</i></code>	Load a module into the current environment
<code>module rm unload <i>modulename</i></code>	Unload a module from the environment
<code>module purge</code>	Unload all loaded modules
<code>module swap <i>module1 module2</i></code>	Swap a loaded module with another
<code>module show <i>modulename</i></code>	Give help for a particular module
<code>module help</code>	Show module specific help

How it works...

Unix shells (Bash, C-shell, etc) use system environment variables to define important values

- By convention, these variables use upper-case letters
- The shell command **printenv** prints all the variables
- To look at a particular one, we use 'echo' and the variable name
 - i.e. to look at the **PATH** variable we type

echo \$PATH

Note the \$ is used to tell shell that we are looking for a variable called PATH

Another important variable is **LD_LIBRARY_PATH** that lists the locations of shared libraries.

Modules work by modifying these variables in a very easy to use manner, independent of the shell you are using.



PATH is used by the shell to search where executables might be. The shell command **which** is used to find the path of an executable.

For example, lets look for where our Python executable may be

```
which python  
/usr/bin/python
```

And we can ask Python which version it is

```
python --version  
Python 2.7.5
```

Suppose we wanted the latest version of Python ?

module avail

So before we install the latest version, or request the Help Desk to do it, let's check the modules

```
module avail python
-----/usr/local/Modules/modulefiles -----
python/2.7.12
```

So let's use it!

```
module load python
```

Let's check we have the right version

```
which python
/usr/local/python/2.7.12/bin/python
python --version
Python 2.7.12
```



Removing Modules

Suppose we wanted to go back to using the system python. This can be done with a number of commands

```
module rm python
```

```
module delete python
```

```
module purge #remove ALL modules, not just python
```

Let's check we have gone back to our original version

```
which python
```

```
/usr/bin/python
```

```
python --version
```

```
module list
```




module Versions

It's possible to have lots of different versions of the same software

```
module avail python
```

```
-----/usr/local/Modules/modulefiles
```

```
python/2.7.12-gcc
```

```
python/2.7.8-gcc(default)
```

```
python/2.7-intel
```

```
python/3.4.3-gcc
```

```
python/3.5.1-gcc
```

Environment module lets you specify the version, or swap between them

module load python #loads the default, same as *module load python/2.7.8-gcc*

which python

```
/usr/local/python/2.7.8-gcc/bin/python
```

module swap python/2.7-intel #unloads the existing Python, then loads the version

```
#compiled with the intel compiler
```

which python

```
/usr/local/python/2.7-intel/bin/python
```



Exercise 6

Build on previous work, write a script (say **interrogate.sh**) to interrogates a node for information

- Print the hostname (*hostname*)
- Print the time of day (*date*)
- Look for an OS system python
 - What is its path?
 - What version is it?
- Load Python via modules
 - What is its path?
 - What version is it?
- Print a farewell message and time

All Good




Need help



- Once you have run this on the login node, write a slurm script to run **interrogate.sh** on a single node
- If you have time, modify the slurm script to run on 2 nodes.
 - Remember to use **srun** when invoking **interrogate.sh**
 - Is the output of the text what you expect?

CHAPTER 7 : Transferring software





scp – Copying files between computers

You can copy files *between* computers with *scp* – ‘secure cp’
scp has similar syntax to *cp*. You specify the username and machine name of the remote host in the command line

scp <from> <to>

where <from> or <to> can also specify a **username@computername:**

e.g.

scp *.db username@computer:/home/apps/dbfiles

Note the ‘:’

The wild card *.db matches any filename whose ending is the string ‘.db’



scp – Copying files between computers

EXAMPLES (Using a computer called **raijin.nci.org.au** with a user called **vader**):

COPYING TWO FILES TO RAIJIN

```
scp mfile.c myfile.h vader@raijin.nci.org.au:
```

COPYING FROM RAIJIN WITH A WILDCARD


```
scp vader@raijin.nci.org.au:mydata/*.dat /home/user
```

Other interesting options ('man cp')

-i <identity_file> use a predefined crypto-key file, removing the need to use a password. Use with caution!

-r copy directories recursively

dos2unix and unix2dos



In Windows (DOS), text files have every line terminated with two characters –CARRIAGE RETURN and LINEFEED =>

Using the C computer language syntax, this maps to `\r\n`

In Unix and Mac, text files have every line terminated with `\n` only

When viewing a DOS file in Unix, you may see a **^M** at the end of every line.
Or worst, the character may appear invisible and cause errors later !

Two tools exist to convert text files between the two Operating systems

dos2unix <filename>

unix2dos <filename>

Making files easier to move

COMPRESSING FILES

Lossless compression of files is useful:

- When transmitting files across the internet
- Saving space in your file system (NCI have quotas)



There are many tools to compress files. One of the most common is **gzip**.

To compress a file:

gzip myfile.txt

Produces a file called myfile.txt.gz (and removes myfile.txt)

To get the file back:

gunzip myfile.txt.gz

Produces a file called myfile.txt (and removes myfile.txt.gz)

Creating and extracting archives

Archiving files

It is often useful to combine a large number of files into one big file

- When transmitting files across the internet
- Quicker access on tape archive systems
- Your file system may have limits on the number of files you can have (N have quotas on this – the inode quota)



There are many tools to archive files. One of the most common is **tar**.

To create an archive:

tar -cf archive.file.name.tar <list of files>

Produces a file called **archive.file.name.tar**

To extract the archive

tar -xf archive.file.name.tar

Extracts all the files in the archive

Once an archive is created, it is a good idea to compress it as well.

This can be done automatically with tar by adding the **-z** flag. i.e. **tar -cfz ...**



Exercise

You want to copy some data from another machine to your home directory. You know the file is called **whatami.tar.gz** and it lives in **/tmp** of a machine called **slurm0**.

Use **scp** to fetch the file. Don't forget the "." at the end of the line.

```
scp slurm0:/tmp/whatami.tar.gz .
```

Check that your the scp worked and you have a file called **whatami.tar.gz** in your directory. Now lets see what is in the archive..

```
tar -xf whatami.tar.gz
```

Use **ls** and **cat** to see what was produced by the file.

All Good



Need help



Chapter 7 Using resources effectively

Although we covered requesting resources from the scheduler earlier, how do we know how much and what type of resources we will need in the first place?

Answer: we don't. Not until we've tried it ourselves at least once. We'll need to benchmark our job and experiment with it before we know how much it needs in the way of resources.

```
sacct -u {userID}
```

```
sacct -j {jobID} -l
```

- **Hostname** - Where did your job run?
- **MaxRSS** - What was the maximum amount of memory used?
- **Elapsed** - How long did the job take?
- **State** - What is the job currently doing/what happened to it?
- **MaxDiskRead** - Amount of data read from disk.
- **MaxDiskWrite** - Amount of data written to disk.



```
[user2@slurm-login ~]$ sacct -u user2
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
54	hostname	all	slurmclass	2	COMPLETED	0:0
55	hostname	all	slurmclass	2	COMPLETED	0:0
56	run.sh	all	slurmclass	1	COMPLETED	0:0
56.batch	batch		slurmclass	1	COMPLETED	0:0
57	mpitest	all	slurmclass	2	COMPLETED	0:0
57.batch	batch		slurmclass	2	COMPLETED	0:0
58	mpitest	all	slurmclass	1	COMPLETED	0:0
58.batch	batch		slurmclass	1	COMPLETED	0:0
59	mpitest	all	slurmclass	3	COMPLETED	0:0
59.batch	batch		slurmclass	3	COMPLETED	0:0



```
[user2@slurm-login ~]$ sacct -j 59
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
59	mpitest	all	slurmclass	3	COMPLETED	0:0
59.batch	batch		slurmclass	3	COMPLETED	0:0

- The state showing that the job is completed successfully.
- The partition that was used in `all`
- The account that was used is `slurmclass`



If you would like to see more information about a job, you can always customise the format for your command

```
sacct -j 59 --format=User,JobID,Jobname,partition,state,time,start,end,elapsed,MaxRss,ncpus,nodelist
```

```
[user2@slurm-login ~]$ sacct -j 59 --format=User,JobID,Jobname,partition,state,time,start,end,elapsed,MaxRss,ncpus,nodelist
```

User	JobID	JobName	Partition	State	Timelimit	Start	End	Elapsed	MaxRSS	NCPUS	NodeList
user2	59	mpitest	all	COMPLETED	05:00:00	2018-08-03T00:35:23	2018-08-03T00:36:19	00:00:56		3	slurm0
	59.batch	batch		COMPLETED		2018-08-03T00:35:23	2018-08-03T00:36:19	00:00:56	32192K	3	slurm0



Why is my job not running?

Use “`squeue --start`” command

Jobs will be listed in order expected start time

Depending upon configuration, not all jobs may have a start time set

Times are only estimates and subject to change



```
01:15:37 m3-login1:~ ctan$ squeue --start
```

JOBID	PARTITION	NAME	USER	ST	START_TIME	NODES	SCHEDNODES	NODELIST(REASON)
2796646	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2796655	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2797577	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2797637	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2798766	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2798776	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2802511	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2802520	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2802526	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2802988	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2802994	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
2809141	comp	GPR_tens	yc1	PD	N/A	1	(null)	(DependencyNeverSatisfied)
3133116	m3i	nt	miguelg	PD	N/A	1	(null)	(launch failed requested held)
3297560	m3g	Refine3D	joshuamh	PD	N/A	1	(null)	(QOSMaxGRESPerUser)
3157584_[464-500]	comp	nCellsDe	mcfadyen	PD	N/A	1	(null)	(QOSMaxCpuPerUserLimit)
3312396	comp	hiplexpi	jste0021	PD	N/A	1	(null)	(Priority)
3312397	comp	hiplexpi	jste0021	PD	N/A	1	(null)	(Priority)
3312365	comp	hiplexpi	jste0021	PD	2018-08-03T01:35:12	1	m3i009	(Resources)
3312385	comp	hiplexpi	jste0021	PD	2018-08-03T01:35:12	1	m3i029	(Priority)
3312393	comp	hiplexpi	jste0021	PD	2018-08-03T01:35:13	1	m3i033	(Priority)
3312394	comp	hiplexpi	jste0021	PD	2018-08-03T01:35:16	1	m3i036	(Priority)
3312395	comp	hiplexpi	jste0021	PD	2018-08-03T01:35:16	1	m3h005	(Priority)
3273020	m3a	hsa_chn2	mziemann	PD	2018-08-03T03:21:01	1	m3a006	(Resources)
3284320	m3a	hsa_chn2	mziemann	PD	2018-08-03T03:21:01	1	m3a007	(Resources)
3130659	comp	oxalipia	sgwe0001	PD	2018-08-03T04:05:56	1	(null)	(Priority)
3221688	m3i	aFe5Dv3A	eyk	PD	2018-08-03T07:10:34	4	(null)	(QOSMaxCpuPerUserLimit)
3223197	m3i	aFe5Dv2B	eyk	PD	2018-08-03T07:10:34	4	(null)	(QOSMaxCpuPerUserLimit)
3073122	comp	238-cpcm	sgwe0001	PD	2018-08-03T07:16:35	1	m3c005	(Priority)
3073134	comp	238-smd-	sgwe0001	PD	2018-08-03T07:16:35	1	(null)	(Priority)
3073164	comp	238-gas-	sgwe0001	PD	2018-08-03T07:16:35	1	(null)	(Priority)
3073192	comp	248-cpcm	sgwe0001	PD	2018-08-03T07:16:35	1	(null)	(Priority)
3073193	comp	248-smd-	sgwe0001	PD	2018-08-03T07:16:35	1	(null)	(Priority)
3073195	comp	248-gas-	sgwe0001	PD	2018-08-03T07:16:35	1	(null)	(Priority)

Be kind to the login node

- The login node is very busy managing lots and lots of jobs! It doesn't have any extra space to run computational work. Don't run jobs on the login node, you can use the interactive session to test out your script.
- Remember, the login node is to be shared with other users.
- If someone is being inappropriate and using the login node to run all of their stuff, message an administrator to take a look at things and deal with them.

Test before scaling

- Before submitting a large run of jobs, submit one as a test first to make sure everything works.

Understanding the filesystem

- always know where to write your data, in the production system, there's usually some space that being backed up and some space might be allocated for scratch
- make sure you check your quota for each of this filesystem



Save time

- Compress files before transferring to save file transfer times with large datasets.
- The less resources you ask for, the faster your jobs will find a slot in which to run. Lots of small jobs generally beat a couple big jobs.

Software tips

- You can generally install software yourself, but if you want a shared installation of some kind, it might be a good idea to message an administrator.
- Always use the default compilers if possible. Newer compilers are great, but older stuff generally has less compatibility issues.



HPC Clusters in Monash University

MASSIVE M3

A Computer for Next-Generation Data Science

1,700 Intel Haswell CPU-cores

2,000 Intel Skylake CPU-cores

48 NVIDIA Tesla K80 GPU

20 NVIDIA Tesla P100 GPU

60 NVIDIA Tesla V100 GPU

coprocessors for data processing and high end visualisation

A 3 petabyte Lustre parallel file system

100 Gb/s Ethernet Mellanox Spectrum

Supplied by Dell, Mellanox and NVIDIA



Monarch

Campus Cluster

Provide Monash researchers with a local capability that focuses on engagement, education and community.

Investment

A co-investor model

Examples include Computational Chemistry, Astro and Fluid Dynamics

1/3rd of MonARCH is co-purchased

Integrated into undergraduate study

CHM3911 Advanced Physical Chemistry

80 students across 3 practical sessions

Gaussian and GaussView for calculations

Students taught how to use a HPC system to perform their calculations



WRAPPING UP

Do you have the tools you need to do your Research?

Would you want to know more?


- Email mcc-help@monash.edu
- Data Fluency Slack Channel
- Bioinformatic Drop-in (every Friday)
- <https://docs.monarch.erc.monash.edu.au>
- <https://docs.massive.org.au>
- HPC web site (i.e. www.nci.org.au)
- Google





Your Feedback please!

ONE THING YOU LIKED



ONE THING TO IMPROVE

